Brainstorming

- Title: Brainstorming — Streamlining Ticket Assignment

- Subtitle: Ideation Phase

Workshop goal

- Generate a broad set of ideas for reducing assignment time, improving assignment accuracy, and balancing agent load.

Structure and rules

- Timebox: 60–90 minutes

- Warm-up: 5 minutes (silent ideation)

- Diverge: 20–30 minutes (quantitative and wild ideas)

- Converge: 20–30 minutes (cluster and vote)

- Criteria: Favor ideas that are measurable, low-risk, and implementable in 3 months

Idea categories

- Automation: rule-based routing, ML-assisted routing, prioritized queues

- Data & Metadata: enrichers, mandatory fields, smart parsers from text

- Agent Experience: skill tags, real-time load dashboards, shift-aware routing

- Process: SLA-based prioritization, auto-escalation rules, human-in-loop checks

- Communication: suggested assignee notifications, canned routing reasons

- Measurement: feedback loop, A/B testing, assignment accuracy telemetry

Seed prompts for ideation

- How might we automatically identify ticket intent with minimal human input?

- What small automations reduce context switches for agents?

- How could we surface agent skills and current load in one glance?

- Which low-cost integrations would give the biggest ROI in 30–90 days?

Generated idea examples (short list)

- Rules-based routing by product tag + priority, fallback to team queue

- Lightweight ML classifier to suggest 3 best assignees (human picks)

- Mandatory "issue category" field with smart suggestions + validation

- Agent load view with predicted occupancy and suggested rebalancing

- Auto-snooze low-priority tickets during peak hours to reduce noise

- Escape hatch: "manual assign" button for exceptions that logs reason

Idea evaluation matrix (simple)

- Columns: Idea Effort (L/M/H) Impact (L/M/H) Risk (L/M/H) Quick wins?

- Example row: Rules-based routing L M-H L Yes

- Use this matrix to prioritize 3 pilots

Prioritization (example outcome)

1. Pilot rules-based routing on product tickets (Quick win)

2. Introduce mandatory category field with smart suggestions

3. Build agent load dashboard + rebalancer in next quarter

4. Prototype ML classifier as human-suggestion engine

Experiment plan for top 2 pilots

- Pilot A: Rules-based routing

    o Duration: 4 weeks

- - Metrics: time-to-assignment, reassignment rate, SLA breaches

  - Success criteria: 30% reduction in time-to-assignment;

- Pilot B: Mandatory category field + smart suggestions

- Duration: 3 weeks (A/B test)

- Metrics: % tickets with valid category; assignment accuracy

- Success criteria: >80% valid category completion; measurable assignment improvement

Next steps and responsibilities

- Action owner: Product led to approve pilot scope

- Engineering: estimate integration work for rules and metadata

- Ops: select pilot teams and define measurement dashboard

- Timeline: kick-off within 2 weeks; pilot results in 6–8 weeks

- Copy each section above into its own document (Docx, Google Doc, or slide)

- Add the cover page and any diagrams (quickly made in Slides or draw.io)

- Export each document as PDF from the app's File > Export or Download as PDF option